

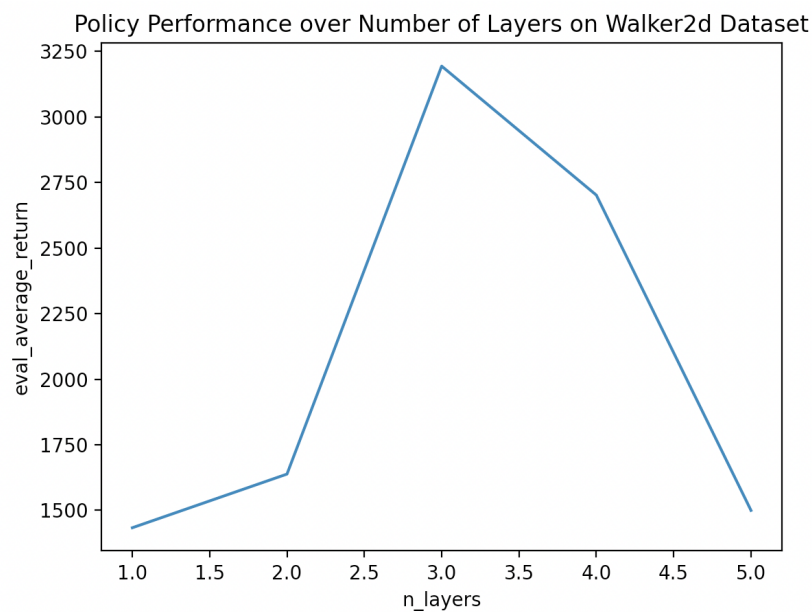
I. Behavioral Cloning

Part 2 See the following table.

Environment	Policy Mean	Policy Std	Expert Mean	% of Expert
Ant	3399.31	102.36	4713.65	72.12%
Walker2d	1638.22	215.35	5566.85	29.43%

Table 1: Comparison of Ant to Walker2d Dataset Using 3000 `eval_batch_size`, 1000 `train_batch_size`, 64 `network_size`, and 1000 `ep_len`

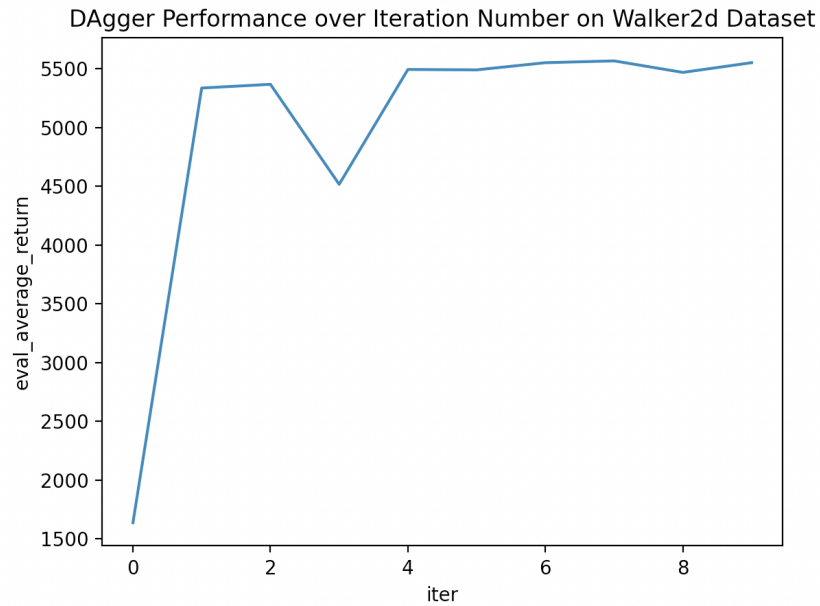
Part 3 See the following graph.



To see the performance change over number of network layers is interesting, because with the same setup (of evaluation batch size, training batch size, network size, etc.), the model seems to underfit for $n < 3$ and overfit for $n > 3$, given by the peak performance at $n = 3$.

II. DAgger

See the following graph for Walker2d.



And the following for Ant.

